

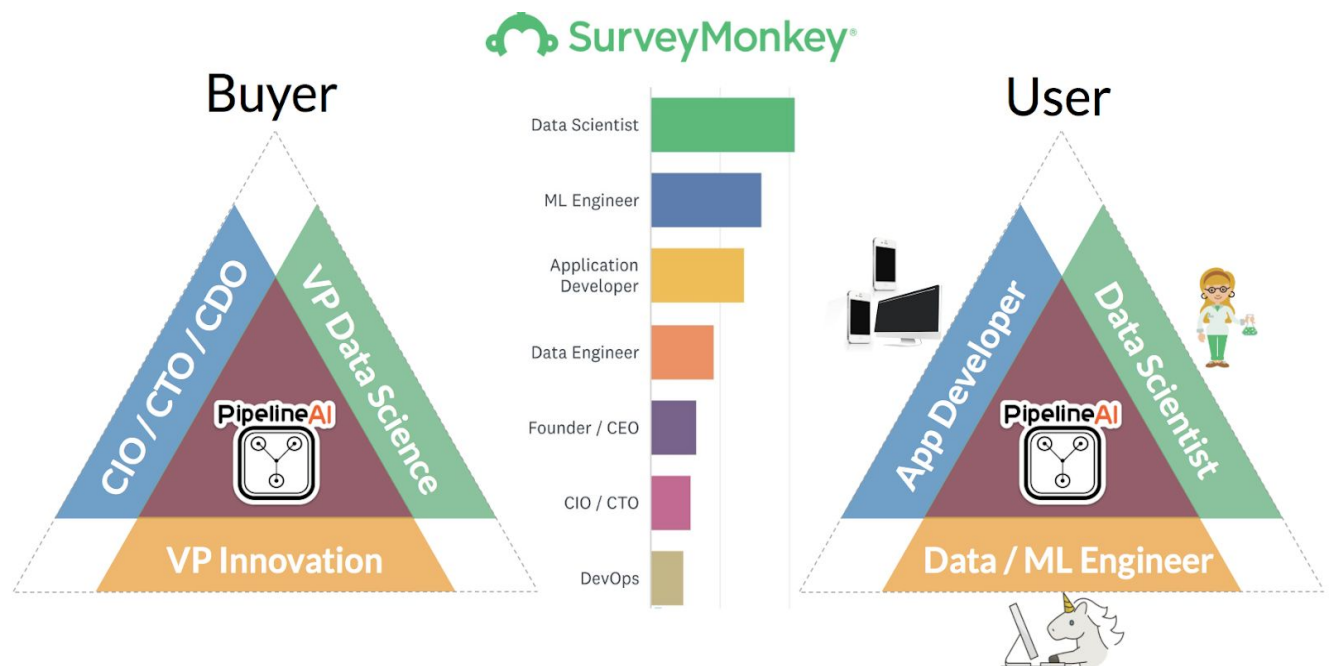
5 Things We Learned From Our Community Survey

1. PipelineAI Satisfies 3 User Personas

We had the opportunity to pose pointed questions via Survey Monkey, initially in September 2018 to our super-engaged 120,000 member community. Enterprise PipelineAI product features stack ranking comes directly from this audience as github issues. We are happy to share the survey insights.

The audience we polled varied from practitioners like Data Scientists, Machine Learning Engineer, Consultants to management and executives. Application Developers, the largest constituents in the developer community unsurprisingly comprise the third highest group.

The typical persona that machine learning platforms like Enterprise PipelineAI appeal to:



PipelineAI: Real-Time Enterprise Machine Learning Platform


<https://pipeline.ai>

5 Things We Learned From Our Community Survey

2. PipelineAI Offers High Business Value

To the question, which aspects of Enterprise PipelineAI brought the most value, the community led with

1. Time to deploy models in production
2. Automated performance increases, auto-tuning
3. Number of models that successfully run in production
4. Concept drift and bias detection
5. API and routing optimizations.

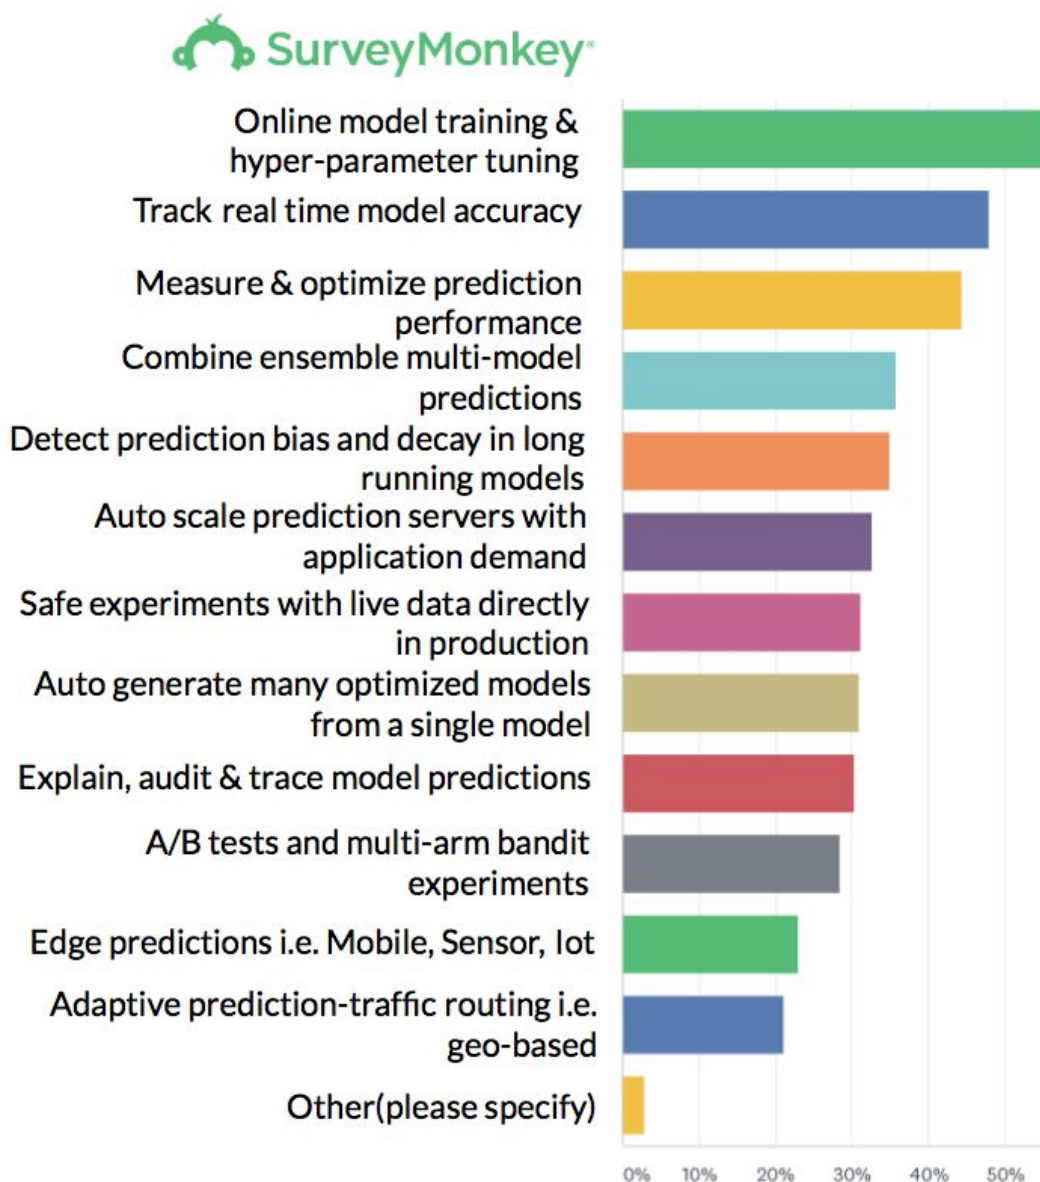


PipelineAI: Real-Time Enterprise Machine Learning Platform

<https://pipeline.ai>


3. PipelineAI is a Feature-Rich Platform

There is no question that online machine learning is a major differentiator. Every business needs online, real-time, machine learning to power their production applications. in inference and therefore in training and hyper-parameter tuning just like Uber, Lyft and Netflix. Top PipelineAI features include the following:



5 Things We Learned From Our Community Survey


4. PipelineAI Runs Anything & Anywhere



It comes as no surprise that Python, Java, R languages, frameworks like Keras, Scikit-Learn, TensorFlow, SparkML and PyTorch have garnered interest from the community. As moving between clouds becomes passe, formats become the norm and we anticipate formats like ONNX will rise up in the list

Azure IBM Watson GCP H2O ML AutoML Databricks AWS IBM
AWS SageMaker Tensorflow Google Cloud
Sagemaker Anaconda Google Azure ML H2O now

Flexibility in deployment is paramount for use adoption. It plays right into sliding machine learning platform infrastructure into the existing CIO blessed guardrails of enterprise architecture. Hybrid and vertical clouds with direct public cloud legs in their own environments is common. Every such scenario needs an orchestration layer that is infrastructure aware. PipelineAI is one such offering.




PipelineAI: Real-Time Enterprise Machine Learning Platform

<https://pipeline.ai>

5. Online ML is Growing in Demand

On the question of how many current and future models and servers for serving, the consensus averaged around 25 models and 120 servers cumulative in 12 months.



Pioneers in machine learning platforms like Uber's Michelangelo confirm adoption hockey sticks past the initial quarter as applications see immediate value. A platform such as PipelineAI is a must have to serve the demand where experimentation in production is at speed of thought. Please email us at sales@pipeline.ai to start using the PipelineAI ML Application Platform.